

Debugging & Profiling

Luca Tornatore, I.N.A.F.

2025 INAF Course on HPC



September, 22nd - 26th, OACT, Catania

Outline




Profiling



Debugging



Profiling outline



Mapping



Profiling
Tools



Profiling
Behaviour



Profiling
Efficiency



Map the workflow

A call tree (more precisely, a **call graph**) is a control flow graph that exhibits the calling relationships among routines in a program.

A *node* in the graph represent a routine, while an *edge* represents a calling relationship.

We'll concentrate on **dynamic call graphs** – i.e. the **records of program executions**.

The most complete graph is context-sensitive, which means that every call stack of a procedure is recorded as a separate node (the resulting graph is called *calling context tree* instead of *call tree*).

However, that requires a larger amount of memory for large program, it is useful in case of code reuse (the same code being executed at different points by different call paths).



Basic profiling concepts

INSTRUMENTATION

Inserts extra code at compile time wrapping function calls to count how many times it calls / is called and how much time it takes to execute.

SAMPLING

The profiler ask for interrupts N_{samples} per sec + interrupts at function calls + interrupts at selected events, and records on a histogram the number of occurrences in every part of the program. The call graph is inferred from these data.

DEBUGGING

The profiler ask for interrupts at every line code and function call (more correct: enters the by-step execution mode)



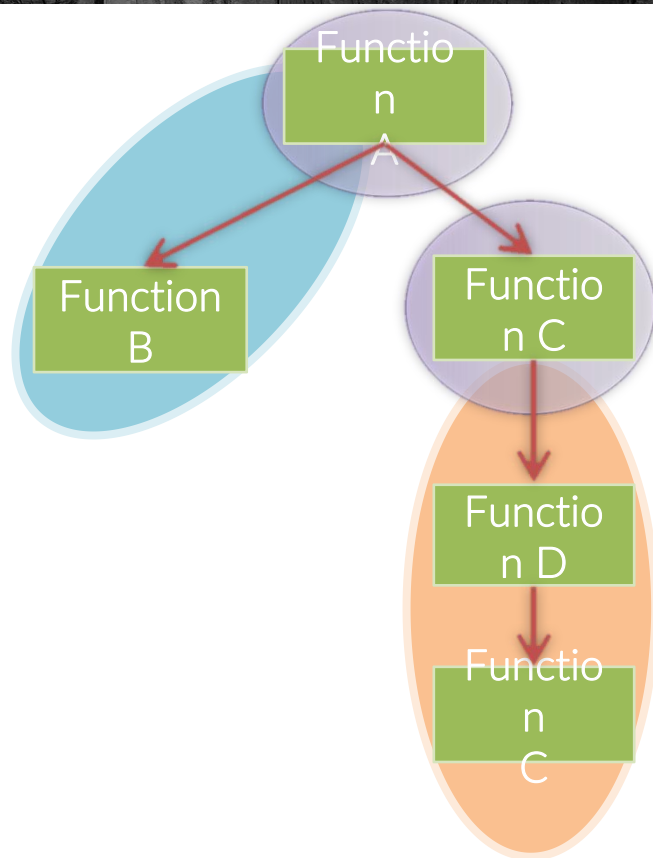
Basic profiling concepts

Collect **program events**

- Hardware interrupts
- Code instrumentation
- Instruction set simulation
- tracing

Periodic sampling

- Top of the stack (exclusive)
- All stack (children inclusive)





How to obtain the call tree

There are several way to obtain a dynamic call tree.

The main open source alternatives are:

1. Using gcc and **gprof**
2. Using **valgrind**
3. Using **perf**
4. Some others, mostly non-free (among significant free: CodeXL by AMD)

Many IDE uses the aforementioned tools or their own plug-ins (eclipse, netbeans, code::block, codelite, ...)



Using gcc + gprof

Just compile your source using **-pg** option.

You should also profile turning on the optimizations you're interested in.

```
gcc -[my.optimizations] -pg -o myprogram.x myprogram.c
```

*Note: don't use the option **-p**. It provides less information than **-pg**.*

After that, run your program normally. Profiling infos will be written in the file `gmon.out`.

▶ You can read (options and details in the man page) the informations by
gprof myprogram.x

▶ You can visualize the call graph by (read the man pages...)

```
gprof myprogram.x | gprof2dot.py | dot -T png -o callgraph.png
```




The call-tree



Well, now that you just met, say goodbye to the glorious **gprof**

He's a dinosaur from the past decades..

- lacks real multithread support
- lacks real line-by-line capability
- does not profile shared libraries
- need recompilation
- may lie easier than other tools (see later..)

You may still consider it for some call counts business

You may still consider it for some call counts business later..)

- may lie easier than other tools (see
- need recompilation



How to induce you profiler to lie

gprof (gcc -pg)
does not record
the call stack

```
#include <stdlib.h>
void loop(int n)
{
    int volatile i; // does not optimize out
    i = 0;
    while(i++ < n);
}

void light(int n) { loop(n); } void heavy(int n)
{ loop(n); }

int main(void)
{
    light(100000);
    heavy(100000000);
    return 0;
}
```



How to induce you profiler to lie

Flat profile:

Each sample counts as 0.01 seconds.

% time	cumulative seconds	self seconds	calls	self ms/call	total ms/call	name
101.30	0.20	0.20	2	101.30	101.30	loop
0.00	0.20	0.00	1	0.00	101.30	heavy
0.00	0.20	0.00	1	0.00	101.30	light



How to induce you profiler to lie

How could it happen?

The man page is clear:

We assume that the time for each execution of a function can be expressed by the total time for the function divided by the number of times the function is called.

Thus the time propagated along the call graph arcs to the function's parents is directly proportional to the number of times that arc is traversed.



How to induce you profiler to lie

VALGRIND seems to better understand the situation, because it doesn't record only the number of calls to a function but also the time spent in a function given a call path (which, at odds, gprof infers backwardly).

That's way the valgrind measure of time spent in a function plus its callees is reliable.

However, it may as well end up to a misleading picture if we stick in and additional layer of complexity, *unless* you explicit tells it to track separately different call stacks with the command-line option - **-separate-callers = N**

Types	Callers	All Callers	Callee Map	Source Code
#	lr	Source		
0		--- From '/home/luca/code/HPC_LECTURES/		
14		void heavy(int n) { loop(n); }		
15				
16		int main(void)		
17	0.00	{		
18	0.00	light(100000);		
	0.10	1 call(s) to 'light' (lie1.g: lie1.c)		
19	0.00	heavy(100000000);		
	99.88	1 call(s) to 'heavy' (lie1.g: lie1.c)		
20	0.00	return 0;		
21	0.00	}		
22				



Using **perf**

[1] COMPILING

```
gcc -g -fno-omit-frame-pointer my_prog.c -o my_prog
```

[2] RUNNING

```
perf record -F ffff -call-graph <fp|lbr|dwarf> \  
my_prog.c <args>
```

[3] ANALYZING

```
perf report --call-graph=graph < --stdio >
```



Using **valgrind**

[1] COMPILING

```
gcc -g -fno-omit-frame-pointer my_prog.c -o my_prog
```

[2] RUNNING

```
Valgrind -tool=callgrind -callgrind-out-file= $CALLGRIND_OUT -dump-  
instr=yes -coollect-jumps=yes -cache-sim=yes -branch-sim=yes < --I1=... >  
< --D1=...> my_prog <args>
```

[3] ANALYZING

```
kcachegrind $CALLGRIND_OUT
```




- **BASIC / SYSTEM tools**

- gprof / gdb / perf / gperftools
- Valgrind – cachegrind, callgrind, ..

- **HARDWARE COUNTER / PMU interface**

- perf
- PAPI
- Intel PMI
- Likwid
- ...

- **HPC Tools**

- HPCtoolkit
- OpenSpeedShop
- TAU
- SCOREP +

- **VENDOR tools**

- ARM-Allinea
- CodeXL (AMD)
- Intel tools
- ...



Tools : valgrind

An instrumentation framework for building dynamic analysis tools.

Valgrind basically runs your code in a virtual “sandbox” where a synthetic CPU (the same you have) is simulated and executes an instrumented code.

There are various Valgrind based tools for debugging and profiling purposes.

- **Memcheck** is a memory error detector → correctness
- **Cachegrind** is a cache and branch-prediction profiler → velocity
- **Callgrind** is a call-graph generating cache profiler. It has some overlap with Cachegrind
- **Helgrind** is a thread error detector → correctness
- **DRD** is also a thread error detector.

Different analysis technique than Helgrind

- **Massif** is a heap profiler → memory efficiency using less memory
- **DHAT** is a different kind of heap profiler → memory layout inefficiencies
- **SGcheck** (experimental tool) that can detect overruns of stack and global arrays

KCacheGrind is a very useful GUI



Memcheck : highlighting memory errors

- Unvalid memory access: overrunning/underrunning of heap blocks or top of stack, addressing freed blocks, ...
- Use of variables with undefined values
- Incorrect freeing of heap memory
- Errors in moving memory (unwanted src/dst overlaps, ...)
- Memory leaks

Cachegrind: simulating the cache

It can report how many hits (L1, L2 and L3, I- and D-) and how many misses.

It can analyze CPU's branch prediction.

`Ir, I1mr, LLmr, Dr, D1mr, DLmr, Bc, Bcm, Bi, Bin, ...`

Callgrind: profiling the CPU

It collects the number of instructions executed, links them to source lines, records the caller/callee relationship between functions, and the numbers of such calls. It can collect data on cache simulation and/or branches.

Outline



Debugging



Debugging outline



Introduction

```
if(q > 0)
{
    switch(q)
    {
        case(1): // row y = 0 and/or plane symmetry
            if(subregions[i][BOTTOM][y] == 0)
            // this subregion in quadrant 1 contains the
            // set-up corners for seed sub-region generation
            SBL[y] = 0; SBL[x] = hmesh - subregion;
            STR[y] = 1; STR[x] = hmesh - subregion;
            // find horizontal extension in this quadrant
            Hp = STR[x] - SBL[x];
            // allocate memory for seeds in this stack
            SEED_y0 = (unsigned int*)malloc(sizeof(unsigned int)*Hp);
            if(!internal_malloc_original_seedtable)
            {
```

Debugging
a code



Inspecting
a code crash



Debugging
running
process(es)



Introducing debugging



Smithsonian



It has long been recognized and documented that insects are the most diverse group of organisms, meaning that the numbers of species of insects are more than any other group. In the world, some 900 thousand different kinds of living insects are known. This representation approximates 80 percent of the world's species. The true figure of living species of insects can only be estimated from present and past studies. Most authorities agree that there are more insect species that have not been described (named by science) than there are insect species that have been previously named. Conservative estimates suggest that this figure is 2 million, but estimates extend to 30 million. In the last decade, much attention has been given to the entomofauna that exists in the canopies of tropical forests of the world. From studies conducted by Terry Erwin of the Smithsonian Institution's Department of Entomology in Latin American forest canopies, the number of living species of insects has been estimated to be 30 million. Insects also probably have the largest biomass of the terrestrial animals. At any time, it is estimated that there are some 10 quintillion (10,000,000,000,000,000) individual insects alive.



Introducing debugging



9/9

0800 Anchan started
1000 " stopped - anchan ✓ { 1.2700 9.037 847 025
1300 (032) MP - MC 1.98264000 9.037 846 995
033) PRO 2 2.130476415 4.6159250
conv 2.130476415
Relays 6-2 in 033 failed special speed test
in relay " 11,000 test.
Relays changed
1100 Started Cosine Tape (Sine check)
1525 Started Multi Adder Test.
1545 Relay #70 Panel F
(moth) in relay.
First actual case of bug being found.
1650 Anchan started.
1700 closed down.

The usage of **bug** and, hence, of **de-bugging** referred to programming is a long-standing tradition, whose origin is difficult to trace back.

An often-told story is about Mark-II calculator located at Harvard: On Sept. 9th, 1945, a technician found a moth in a relay that caused a flaw in a “program” execution.

Adm. Grace Hopper is reported to have written in its diary about that as “first *actual* case of bug being found”.



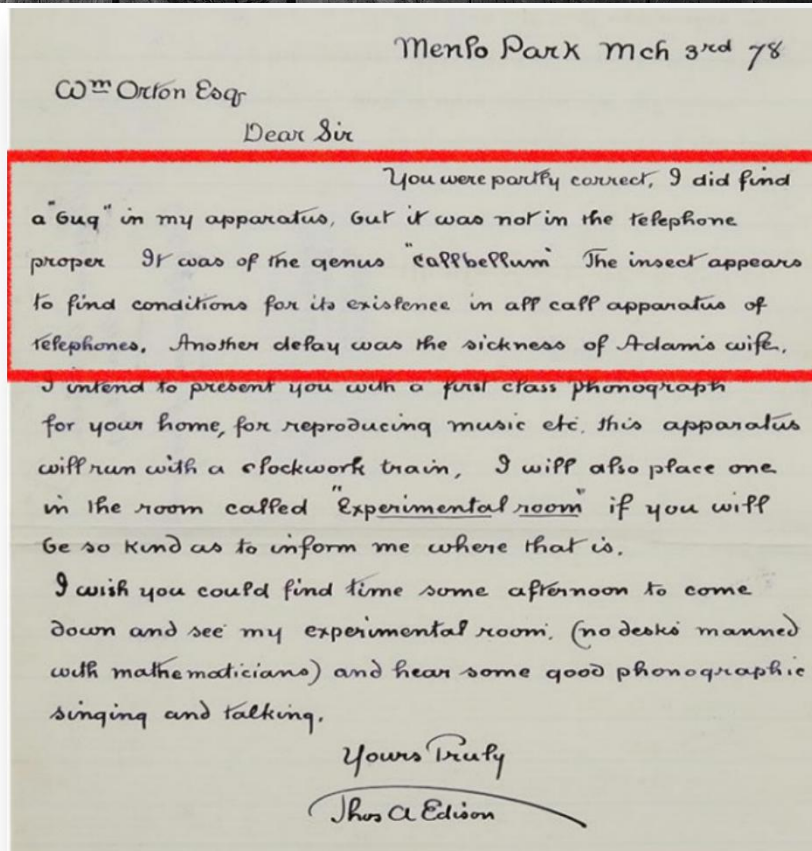
Introducing debugging



However, **Thomas Edison** found another *actual* bug in one of his phones, as he reports in a letter to an associate.

He later writes:

"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that "**bugs**"—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."





Introducing debugging



It is difficult to trace back how this term has been poured into computer programming jargon. However, already in the early 60s it was appearing in technical papers without need of explanation.

The immortal I. Asimov used it in a 1944 short robot story “Catch that rabbit”, and his enormous influence contributed much to make the term popular.

More funny infos:

[IEEE Annals of the History of Computing](#)
(Volume: 20 , [Issue: 4](#) , Oct-Dec 1998)

ask for the PDF if you're interested and do not have access

Stalking the Elusive Computer Bug

PEGGY ALDRICH KIDWELL

From at least the time of Thomas Edison, U.S. engineers have used the word “bug” to refer to flaws in the systems they developed. This short word conveniently covered a multitude of possible problems. It also suggested that difficulties were small and could be easily corrected. IBM engineers who installed the ASSC Mark I at Harvard University in 1944 taught the phrase to the staff there. Grace Murray Hopper used the word with particular enthusiasm in documents relating to her work. In 1947, when technicians building the Mark II computer at Harvard discovered a moth in one of the relays, they saved it as the first actual case of a bug being found. In the early 1950s, the terms “bug” and “debug,” as applied to computers and computer programs, began to appear not only in computer documentation but even in the popular press.

Introduction

Talking computer bugs—that is to say, finding errors in computer hardware and software—occupies and has occupied much of the time and ingenuity of the people who design, build, program, and use computers.¹ Early programmers realized this with some distress. Maurice Wilkes recalls that in about June of 1949:

I was trying to get working my first non-trivial program, which was one for the numerical integration of Airy's dif-

published in the *Annals* in 1981. Here the time is given as the summer of 1945, but the computer is the Mark II, not the Mark I. A photograph shows the moth taped in the logbook, labeled “first actual case of bug being found.” Small problems with computers have been called bugs ever since.⁵

6/4



Introducing debugging



That maybe is too dramatic and emphatic, but the point to get from Edison is that the *de-bugging activity* is an inherent and intrinsic one in software development

*"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that **"bugs"**—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."*

T. Edison



Introducing debugging



Provided that

- you'll have close encounters with bugs in your life;
- de-bugging is a fundamental and unavoidable part of your work;

what is the best way to proceed ?



Introducing debugging



1. Do not insert bugs

Highly encouraged, but rarely works



1. Do not insert bugs

Highly encouraged, but rarely works

2. Add `printf` statements everywhere

Highly discouraged, but sometimes it works.

In case of memory problems – typically due to pointers chaos – you may see that the problem “disappear”, or changes its appearance, when you insert a new `printf`



Introducing debugging



1. Do not insert bugs

Highly encouraged, but rarely works

2. Add `printf` statements everywhere

Highly discouraged, but sometimes works.

In case of memory problems – typically due to oops

3. Use a DEBUGGER

That is definitely the best choice and, fortunately, the subject of this lecture.



Introducing debugging



`gdb` is almost certainly the best free,
extremely feature-rich command-line
debugger ubiquitously available on
*nix systems.



There are 3 basic usages^(*) of GDB:

1. Debugging a code
quite mandatory to compile with -g
2. Inspecting a code crash through a core file
3. Debugging / inspecting a running code

(*) Highly advised: learn keyboard commands. Although many gui exist (we'll see some later), keyboard is still the best productivity tool.




Debugging outline



Introduction

```
if(q > 0)
{
    switch(q)
    {
        case(1): // row y = 0 and/or plane symmetry
            if(subregions[i][BOTTOM][y] == 0)
            {
                // this subregion in quadrant 1 contains the
                // set-up corners for seed sub-region generation
                SBL[y] = 0; SBL[x] = hmesh - subregion;
                STR[y] = 1; STR[x] = hmesh - subregion;
                // find horizontal extension in this quadrant
                Hp = STR[x] - SBL[x];
                // allocate memory for seeds in this stack
                SEED_y0 = (unsigned int*)malloc(sizeof(unsigned int)*Hp);
                if(!internal_malloc_original_seedtable)
                {
```

Debugging
a code



Inspecting
a code crash



Debugging
a running
process(es)



Compiling with dbg infos



In order to include debugging information in your code, you need to compile it with `-g` family options (read the gcc manual for complete info):

`-g`

produce dbg info in O.S. native format

`-ggdb`

produce gdb specific extended info, as much as possible

`-glevel`

default level is 2. 0 amounts to no info, 1 is minimal, 3 includes extra information (for instance, macros expansion) – this allows macro expansion; add **`-gdwarf-n`** in case, where possible, where *n* is the maximum allowed (4)

`-ggdblevel`

you can combine the two to maximize the amount of useful info generated

remember: **`-fno-omit-frame-pointer`**, especially if you are using `-Ox`



Debugging: start



You just start your code under gdb control:

```
%> gdb program
```

You can define the arguments needed by your program already at invocation:

► **live demo with** `gdb_try_args.c`

```
%> gdb --args program arg1 arg2 ... argN
```

Or you can define the arguments from within the gdb session:

```
%> gdb program
Reading symbols from program...done.
(gdb) set args arg1 arg2 ... argN
(gdb) run
```



Debugging
a code

Debugging: run & stop



You may just want the code to run, for instance to reach the point of a seg fault:

```
%> (gdb) run
```

Or, you may want to stop it from the beginning to have full control of each step:

```
%> (gdb) break main
```

```
%> (gdb) run
```

Or you may already know what is the problematic point to stop at:

```
%> (gdb) break location
```

```
%> (gdb) run
```



Debugging: run & stop



Breakpoints are a key concept in debugging. They are stopping point at which the execution interrupts and the control is given back to you, so that you can inspect the memory contents (variables values, registers values, ...) or follow the subsequent execution step by step.

You can define a breakpoint in several way

(gdb) `break`

insert a break at the current pos

(gdb) `break ±
offset`

insert a break *offset* lines
after/before the current line

(gdb) `break
filename:linenum`

insert a break at *linenum* of file
filename

(gdb) `break
functionname`

insert a break at the entry point of
function *functionname*

There are more
options, just
check the manual



Debugging: run & stop



A breakpoint may be defined as dependent on a given condition:

```
(gdb) break my_function if (arg1 > 3 )
```

This sets a breakpoint at function *my_function* : the condition will be evaluated each time the point is reached, and the execution is stopped only if it is true.

Condition can be any valid expression.

```
(gdb) info break
```

gives you informations on active breakpoints.

```
(gdb) delete [n]  
clear [location]  
<disable | enable> see the manual
```



Debugging: run & stop



You can define a list of commands to be executed when a given breakpoint is reached:

```
(gdb) break my_function if (arg1 > 3 )  
Breakpoint 1 at 0x.....: file blabla.c, line 42  
(gdb) command 1
```

Type commands for when breakpoint 1 is hit, one per line.

End with a line saying just "end".

```
> print arg1  
> print another_useful_variable  
> x/10wd a_global_integer_array  
end
```



Debugging: run & stop



When you have the control of the program execution, you can decide how to proceed:

(gdb) `cont [c]`

continue until the end / next stop

(gdb) `cont count-ignore`

continue ignoring the next *count-ignore* stops (for instance, a bp)

(gdb) `next [n] | count
nexti`

continue to the next src line *in the current stack frame*

(gdb) `step [s] | count
stepi`

continue to the next src line

(gdb) `until [u] | count`

continue until a src line past the current one is reached in the current stack fr.

(gdb) `advance location`

continue until the specified location is reached



Debugging: run & stop



You can also rewind your execution step-by-step

```
(gdb) reverse-continue [rc]
```

```
(gdb) reverse-step [count]  
reverse-stepi
```

```
(gdb) reverse-next [count]  
reverse-nexti
```

```
(gdb) set exec-direction <verse | forward >
```



Debugging: source list



Often, when you are debugging, you may have the need of looking at either the source lines or at the generated assembler:

- | | |
|--|--|
| <code>(gdb) list <i>linenum</i></code> | print src lines around line
<i>linenum</i> of the current source file |
| <code>(gdb) list <i>function</i></code> | print the source lines of <i>function</i> |
| <code>(gdb) list <i>location</i></code> | print src lines around <i>location</i> |
| <code>(gdb) set listsize
<i>count</i></code> | control the number of src lines
printed |
| <code>(gdb) disass
[/m] [<i>function</i>] [<i>locati</i>
<i>on</i>]</code> | show the assembler |



Debugging: stack examination



Examining the stack is often of vital importance. With GDB you can have a quick and detailed inspection of all the stack frames.

<code>backtrace</code>	print the backtrace of the whole stack
<code>backtrace <i>n</i></code>	print only the <i>n</i> innermost frames
<code>backtrace -<i>n</i></code>	print only the <i>n</i> outermost frames
<code>backtrace full</code>	print local variables value, also

where, `info stack` additional aliases



Accessing to the content of memory is a fundamental ability of a debugger. You have several different ways to do that:

```
(gdb) print variable  
p/F variable
```

print the value of *variable*
print *variable* in a different format
(*x, d, u, o, t, a, c, f, s*)
see then manual for advanced location

```
(gdb) x/FMT address
```

Explore memory starting at address *address*
-> see at live demo how to use this

```
(gdb) display expr
```

Add *expr* to the list of expressions to display each time your program stops

```
display/fmt expr  
display/fmt addr
```



Memory can be searched to find a particular value,
of a given size

```
(gdb) find [/sn]  
start, end, val1  
[,val2, ...]
```

```
find [/sn]  
start, +len, val1  
[,val2, ...]
```

Search memory for a
particular sequence of
bytes.

s is the size of type to be
searched

n is the max number of
occurrences



Debugging: registers examination



Examining registers may be also useful (although it's something that only quite advanced users can conceive)

- | | |
|-----------------------------------|--|
| <code>(gdb) info registers</code> | print the value of all registers |
| <code>(gdb) info vector</code> | print the content of vector registers |
| <code>(gdb) print \$rsp</code> | print value of the stack pointer |
| <code>(gdb) x/10wd \$rsp</code> | print values of the first 10 4-bytes integers on the stack |
| <code>(gdb) x/10i \$rip</code> | print the next 10 asm instructions |



If there are macros in your code, they can be expanded, provided that you compiled the code with the appropriate option:

`-g3 [gdb3] [-dwarf-4]`

`(gdb) macro expand macro`

shows the expansion of macro *macro*;
expression can be any string of tokens

`(gdb) info macro [-a|-all]
macro`

shows the current (or all) definition(s) of
macro

`(gdb) info macros location`

shows all macro definitions effective at
location



Debugging: watch points



You can set *watchpoints* (aka “keep an eye on this and that”) instead of breakpoints, to stop the execution whenever a value of an expression / variable / memory region changes

(gdb) watch *variable*

keep an eye onto *variable*

(gdb) watch *expression*

stops when the value of expression changes (the scope of variables is respected)

(gdb) watch -l *expression*

Interpret *expression* as a memory location to be watched

(gdb) watch -l *expression*
[mask *maskvalue*]

a mask for memory watching:
specifies what bits of an address should be ignored (to match more addresses)

(gdb) rwatch [-l] *expr*
[mask *mvalue*]

stops when the value of *expr* is read



Debugging outline



Inspecting
a code crash



Debugging
a running
process(es)



After a crash : having the core file



It happens that you have code crashes in conditions not easily reproducible when you debug the code itself, for a number of reasons.

However, the O.S. can dump the entire “program status” on a file, called the *core file*:

```
luca@GGG:~/code/tricks% ./gdb_try_watch
no arguments were given, using default: 100

something wrong at point 2
[1] 8435 segmentation fault (core dumped) ./gdb_try_watch
luca@GGG:~/code/tricks% ls -l core
-rw----- 1 luca luca 413696 nov  8 15:45 core
luca@GGG:~/code/tricks% □
```

In order to allow it to dump the core, you have to check / set the core file size limit:

```
%> ulimit -c [size limit in KB]
```



Once you have a core, you can inspect it with GDB

```
%> gdb executable_name core_file_name
```

Or

```
%> gdb executable_name  
(gdb) core core_file_name
```

The first thing to do, normally, is to unwind the stack frame to understand where the program crashed:

```
(gdb) bt full
```



Debugging outline



Introduction



Debugging
a code



Inspecting
a code crash



Debugging
a running
process



Attach to a process



In order to debug a running process, you can simply attach gdb to it:

```
%> gdb  
(gdb) attach process-id
```

and start searching it to understand what is going on



Debugging outline



Debugging
parallel
processes





Debugging in parallel

The problem is much more complex: the fundamental additional challenge is the simultaneous execution.

Shared memory paradigm: OpenMP, pthreads, ...

- Multiple threads running
- Shared vs private memory regions
- Race conditions

Message-passing paradigm: MPI

- Multiple independent processes (+ possible multithread)
- Communication
- Deadlocks



Multi-threads capability of gdb

Let's have a try:

```
:~$ gcc -g -o my_threadprog my_threadprog.c -lpthread
```

```
:~$ gdb ./my_threadprog
```



Multi-threads capability of gdb

```
:~$ gdb ./my_threadprog
Reading symbols from my_threadprog...done.

(gdb) r

(gdb) Starting program: my_threadprog

[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".
Creating thread 0
[New Thread 0x7ffff77ef700 (LWP 24144)]
Thread #0 says: " Hello World! "
Creating thread 1
[New Thread 0x7ffff6fee700 (LWP 24145)]
Thread #1 says: " Hello World! "
Creating thread 2
[New Thread 0x7ffffffffff700 (LWP 24146)]
Thread #2 says: " Hello World! "
Creating thread 3
[New Thread 0x7ffff65d6700 (LWP 24147)]
Thread #3 says: " Hello World! "
[Thread 0x7ffffffffff700 (LWP 24146) exited]
[Thread 0x7ffff6fee700 (LWP 24145) exited]
[Thread 0x7ffff77ef700 (LWP 24144) exited]
[Thread 0x7ffff7fac700 (LWP 24140) exited]
[Inferior 1 (process 24140) exited normally]
(gdb) Quit
(gdb)
```



Multi-threads capability of gdb

It is necessary to explicitly set up **gdb** for multi-thread debugging

```
(gdb) set pagination off
```

```
(gdb) set target-async on
```

```
(gdb) set non-stop on
```



Multi-threads capability of gdb

In **all-stop** mode, whenever the execution stops, *all* the threads stop (wherever they are).

Whenever you restart the execution, *all* the threads re-start: however, **gdb** can not single-step all the threads in the steplock. Some threads may execute several instructions even if you single-stepped the thread under focus with *step* or *next* commands.

non-stop mode means that when you stop a thread, all the other ones continue running until they finish or they reach some breakpoint that you pre-defined

→ live demo



Multi-threads capability of gdb

Change focus to *thread_no*

Shows info on active threads.
* tags the active thread

(gdb) thread *thread_no*

(gdb) info threads

(gdb) thread apply [*thread_no*] [*all*] *args*

(gdb) break <...> thread *thread_no*

Insert a break into a list
of threads

Apply a command to a list
of threads



Some hint on the workflow

- ▶ If possible, write a code natively parallel but able to run in serial, which means with 1 MPI task or 1 thread
- ▶ Profile, debug and optimize that code in serial first
- ▶ If multi-threaded, deal with thread sync / races with 1 MPI task



Some hint on the workflow

- ▶ Deal with communications, synchronization and race/deadlock conditions on a *small* number of MPI tasks
- ▶ Profile, debug and optimize communications on a *small* number of MPI tasks
- ▶ Get to the full-size run
unfortunately, some times bugs or improper design issues arise only with large number of processes or threads



gdb and MPI

It is still possible to use **gdb** directly, called from mpirun:

```
:~$ mpirun -np <NP> -e gdb ./program
```

However, depending on your system that may not work properly.



| gdb and MPI

The simplest way to use **gdb** with a parallel program is :

```
:~$ mpirun -np <NP> xterm -hold -e gdb ./program
```

Which launches <NP> xterm windows with running gdb processes in which you can run each parallel process

```
(gdb) run <arg_1> <arg_2> ... <arg_n>
```

or

```
... xterm -hold -e gdb -args ./program <arg_1> ...
```



gdb and MPI

```
:~$ mpirun -np <NP> xterm -hold -e gdb --args ./program <arg_1> ... <arg_2>
```

On a HPC facility, normally you do that while running an *interactive session*..
..and in several occasion *this* will not work, because HPC environments are hostile to X for several reasons (remember to connect with `-X` or `-Y` switch of ssh).



Another not so handy possibility is to open as many connections as processes on different terminals on your local machine, and *attach* **gdb** to the already running MPI processes

```
:~$ mpirun -np <NP> ./program
```

Followed by:

```
:~$ gdb -p <PID_of_MPI_task_n>
```

For each MPI task you want to follow.



There are still 2 issues

1. *Where* to run **gdb**, if **xterm** is not available and you do not want to use it in multi-thread mode ?

You may consider using **screen** (practical example in few minutes)

2. *How* the MPI tasks should be convinced to wait for **gdb** to step in ?

→ *Next slides*



Note

A possible issue for attaching gdb to a running process is that you may not have the capability to do that on a Linux system.

Look in the file:
`/proc/sys/kernel/yama/ptrace_scope`

- 0 ("classic ptrace permissions")
No additional restrictions on operations that perform **PTRACE_MODE_ATTACH** checks (beyond those imposed by the commoncap and other LSMs).
The use of **PTRACE_TRACEME** is unchanged.
- 1 ("restricted ptrace") [default value]
When performing an operation that requires a **PTRACE_MODE_ATTACH** check, the calling process must either have the **CAP_SYS_PTRACE** capability in the user namespace of the target process or it must have a predefined relationship with the target process. By default, the predefined relationship is that the target process must be a descendant of the caller.
A target process can employ the `prctl(2)` **PR_SET_PTRACER** operation to declare an additional PID that is allowed to perform **PTRACE_MODE_ATTACH** operations on the target. See the kernel source file [Documentation/admin-guide/LSM/Yama.rst](#) (or [Documentation/security/Yama.txt](#) before Linux 4.13) for further details.
The use of **PTRACE_TRACEME** is unchanged.
- 2 ("admin-only attach")
Only processes with the **CAP_SYS_PTRACE** capability in the user namespace of the target process may perform **PTRACE_MODE_ATTACH** operations or trace children that employ **PTRACE_TRACEME**.
- 3 ("no attach")
No process may perform **PTRACE_MODE_ATTACH** operations or trace children that employ **PTRACE_TRACEME**.



gdb and MPI

We are left with the problem of *attaching* the **gdb** to a running process (or several running processes).

There is a classical trick, that requires to insert some small additional code in your program

```
int wait = 1
```

```
while(wait)  
    sleep(1);
```

The MPI processes will wait indefinitely until the value of wait does not change..
*which you can do from inside **gdb** attached to each process.*



Note

Solutions:

1. Get the capability
2. As root type:
`echo 0 > /proc/sys/kernel/yama/ptrace_scope`
3. Set the `kernel.yama.ptrace_scope` variable in the file `/etc/sysctl.d/10-ptrace.conf` to 0

The last solution turns off the security measure permanently, it is not a good idea (at least on a facility)

- 0 ("classic ptrace permissions")
No additional restrictions on operations that perform **PTRACE_MODE_ATTACH** checks (beyond those imposed by the commoncap and other LSMs).
The use of **PTRACE_TRACEME** is unchanged.
- 1 ("restricted ptrace") [default value]
When performing an operation that requires a **PTRACE_MODE_ATTACH** check, the calling process must either have the **CAP_SYS_PTRACE** capability in the user namespace of the target process or it must have a predefined relationship with the target process. By default, the predefined relationship is that the target process must be a descendant of the caller.
A target process can employ the `prctl(2) PR_SET_PTRACER` operation to declare an additional PID that is allowed to perform **PTRACE_MODE_ATTACH** operations on the target. See the kernel source file [Documentation/admin-guide/LSM/Yama.rst](#) (or [Documentation/security/Yama.txt](#) before Linux 4.13) for further details.
The use of **PTRACE_TRACEME** is unchanged.
- 2 ("admin-only attach")
Only processes with the **CAP_SYS_PTRACE** capability in the user namespace of the target process may perform **PTRACE_MODE_ATTACH** operations or trace children that employ **PTRACE_TRACEME**.
- 3 ("no attach")
No process may perform **PTRACE_MODE_ATTACH** operations or trace children that employ **PTRACE_TRACEME**.



gdb and MPI

Let's say that your MPI program starts with:

```
int main(int argc, char **argv)
{
    int Me, Size;

    MPI_Init(&argc, &argv);

    MPI_Comm_rank(MPI_COMM_WORLD, &Me);
    MPI_Comm_size(MPI_COMM_WORLD, &Size);

    ...
}
```

and that you insert the following code snippets right after it →



gdb and MPI

compiled only if DEBUGGER is defined

```
#ifdef DEBUGGER
    int wait = 1;
    pid_t my_pid;
    char my_host_name[200];

    gethostname(my_host_name, 200);
    my_pid = getpid();
```

```
    for(int i = 0; i < Size; i++)
    {
        if(i == Me)
            printf("task with PID %d on host %s is waiting\n",
                  my_pid, my_host_name);
        MPI_Barrier(MPI_COMM_WORLD);
    }
```

```
    while ( wait )
        sleep(1);
```

} wait is 1, so each process is just spinning

```
    MPI_Barrier(MPI_COMM_WORLD);
```

```
#endif
```

} Once MPI procs exit the previous while they are not rushing away

Each process print the
message, forced to follow
rank-order



A little bit more of flexibility:

```
char *env_ptr;  
if( ( env_ptr = getenv("DEBUG_THIS")) != NULL) &&  
    ( strncasecmp(env_ptr, "YES", 3) == 0 )  
{  
    int    wait = 1;  
    pid_t my_pid;  
    char   my_host_name[200];  
  
    gethostname(my_host_name, 200);  
    my_pid = getpid();  
  
    < ... >  
  
    while ( wait )  
        sleep(1);  
  
    MPI_Barrier(MPI_COMM_WORLD);  
}
```

Now this code is always there.

It becomes active only when you define the environment variable “DEBUG_THIS” as being “YES”, which you can do at the shell prompt right before calling mpirun.



Even more flexibility:

```
char *env_ptr;  
if( ( env_ptr = getenv("DEBUG_THIS")) != NULL) &&  
    ( strncasecmp(env_ptr, "YES", 3) == 0 )  
{  
    int    get_through = Me;  
    pid_t my_pid;  
    char   my_host_name[200];  
  
    gethostname(my_host_name, 200);  
    my_pid = getpid();  
  
    < ... >  
  
    while ( get_through )  
        sleep(1);  
  
    MPI_Barrier(MPI_COMM_WORLD);  
}
```

get_through is set to the MPI rank of the process, i.e. it is > 0 for all the procs but for the 0 one.

In other words, all the MPI tasks but the 0th will wait at the subsequent barrier.

This way, you can avoid to manually change **get_through** for *all* the tasks and unlock only the 0th



gdb and MPI



A handy alternative to is to use **screen** on a single term.
screen is a utility you may have to install by yourself (or ask the sys admin to do that):
check your preferred packaging system or whatever you use to install apps (on HPC facilities you shall compile and install in your home, however).
Basic commands:

<code>screen</code>	start the screen
<code>ctrl+a ?</code>	get help on commands
<code>“</code>	get the list of active windows
<code>c</code>	create a new window
<code>A</code>	rename the current window

Debugging outline



GUI
for GDB



1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS
5. DDD
6. NEMIVER, ECLIPSE, NETBEANS, CODEBLOCKS,
many others..



You can start gdb with a text-user-interface:

```
%> gdb -tui
```

Or you can activate/deactivate it from gdb itself:

```
(gdb) ctrl-x a
```

```
(gdb) ctrl-x o
```

```
(gdb) ctrl-x 2
```

```
(gdb) layout src
```

```
asm
```

```
split
```

```
regs
```

change focus

shows assembly windows

display src and commands

display assembly and commands

display src, asm and commands

display registers window



GDB built-in tui



```
gdb_try_breaks.c
372 {
373
B+> 374 if ( argc > 1)
375     // arg 0 is the name of the program itself
376     {
377         printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378         for ( int i = 1; i < argc; i++ )
379             {
380                 printf( "\targument %d is : %s\n", i, *(argv+i) );
381             }
382         printf( "\n" );
383     }
384
385     else
386
387         printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );
388
389
390     int arg1;
391
392     if ( argc > 1 )
393         arg1 = atoi( *(argv+1) );
394
395     else
396         arg1 = DEFAULT_ARG1;
397
398     int ret;
399
400     ret = function 1( arg1 );

native process 8943 In: main                                     L374  PC: 0x555555554d38
(gdb) l
360     in /home/luca/code/tricks/gdb_try_breaks.c
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks

Breakpoint 1, main (argc=1, argv=0x7fffffffdaa8) at gdb_try_breaks.c:374
(gdb) □
```



GDB built-in tui



```

gdb_try_breaks.c
B+> 374 if ( argc > 1)
375     // arg 0 is the name of the program itself
376     {
377         printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378         for ( int i = 1; i < argc; i++ )
379             {
380                 printf( "\targument %d is : %s\n", i, *(argv+i) );
381             }
382         printf( "\n" );
383     }
384
385     else
386
387         printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );

B+> 0x55555554d38 <main+15>      cmpl    $0x1,-0x14(%rbp)
0x55555554d3c <main+19>      jle     0x55555554db7 <main+142>
0x55555554d3e <main+21>      cmpl    $0x2,-0x14(%rbp)
0x55555554d42 <main+25>      jle     0x55555554d4b <main+34>
0x55555554d44 <main+27>      mov     $0x73,%edx
0x55555554d49 <main+32>      jmp     0x55555554d50 <main+39>
0x55555554d4b <main+34>      mov     $0x20,%edx
0x55555554d50 <main+39>      mov     -0x14(%rbp),%eax
0x55555554d53 <main+42>      sub     $0x1,%eax
0x55555554d56 <main+45>      mov     %eax,%esi
0x55555554d58 <main+47>      lea     0x2bf(%rip),%rdi      # 0x5555555501e
0x55555554d5f <main+54>      mov     $0x0,%eax
0x55555554d64 <main+59>      callq   0x55555554d680 <printf@plt>
0x55555554d69 <main+64>      movl    $0x1,-0xc(%rbp)

native process 9102 In: main                                     L374  PC: 0x55555554d38
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks

Breakpoint 1, main (argc=1, argv=0x7fffffffdaa8) at gdb_try_breaks.c:374
(gdb) layout split

```



GUI for
GDB

GDB dashboard



<https://github.com/cyrus-and/gdb-dashboard>

```
dash
Source
1 void fun(int n, char *data[])
2 {
3     int i;
4     for (i = 0; i < n; i++) {
5         printf("%d: %s\n", i, data[i]);
6     }
7 }
8
9 int main(int argc, char *argv[])
10 {
11     int i;
12     for (i = 0; i < argc; i++) {
13         fun(i, argv[i]);
14     }
15 }
16
17 Assembly
18 0x0000000010000f18 b0 00 fun+56 mov $0x0,%al
19 0x0000000010000f1a e8 4b 00 00 00 fun+58 callq 0x100000f6a
20 0x0000000010000f1f 89 45 e8 fun+63 mov %eax,-0x18(%rbp)
21 0x0000000010000f22 8b 45 ec fun+66 mov -0x14(%rbp),%eax
22 0x0000000010000f25 85 01 00 00 fun+69 add $0x1,%eax
23 0x0000000010000f2a 89 45 ec fun+74 mov %eax,-0x14(%rbp)
24 0x0000000010000f2d e9 c4 ff ff fun+77 jmpq 0x100000ef6 <fun+22>
25 0x0000000010000f32 48 83 c4 20 fun+82 add $0x20,%rsp
26 0x0000000010000f36 5d fun+86 pop %rbp
27 0x0000000010000f37 c3 fun+87 retq
28
29 Threads
30 [1] id 4355 from 0x0000000010000f2a in fun+74 at scrot.c:7
31
32 Stack
33 [0] from 0x0000000010000f2a in fun+74 at scrot.c:7
34 arg n = 3
35 arg data = 0x7fff5bffb60
36 loc i = 1
37 [1] from 0x0000000010000f62 in main+34 at scrot.c:14
38 arg argc = 3
39 arg argv = 0x7fff5bffb60
40 (no locals)
41
42 Registers
43 rax 0x0000000000000002 rbx 0x0000000000000000 rcx 0x0000010000000203
44 rdx 0x0000020000000200 rsi 0x0000000000012068 rdi 0x00007fff79e86118
45 rbp 0x00007fff5bffb200 rsp 0x00007fff5bffb000 r8 0x0000000000000040
46 r9 0x00007fff79e86110 r10 0xffffffffffffffff r11 0x0000000000000246
47 r12 0x0000000000000000 r13 0x0000000000000000 r14 0x0000000000000000
48 r15 0x0000000000000000 rip 0x0000000100000f2a eflags [ TF IF ]
49 cs 0x0000002b ss <unavailable> ds <unavailable>
50 es <unavailable> fs 0x00000000 gs 0x00000000
51
52 Expressions
53 [1] data[i] = 0x7fff5bffb60 "hello"
54
55 Memory
56 0x00007fff5bffb0c 01 00 00 00 60 fb bf 5f ff 7f 00 00 00 00 00 00 .....b.....
57 0x00007fff5bffb1c 03 00 00 00 40 fb bf 5f ff 7f 00 00 62 0f 00 00 .....b.....
58 0x00007fff5bffb2c 01 00 00 00 60 fb bf 5f ff 7f 00 00 03 00 00 00 .....b.....
59 0x00007fff5bffb3c 00 00 00 00 50 fb bf 5f ff 7f 00 00 ad 15 f7 9c .....P.....
60
61 0x00007fff5bffb60 68 65 6c 6c 6f 00 47 44 42 00 54 45 52 4d 5f 50 hello.GDB.TERM_P
62
63 History
64 $$1 = 63
65 $$0 = {[0] = 0x7fff5bffb60 "hello", [1] = 0x7fff5bffb60 "GDB"}
```



GUI for
GDB

GDBgui



[gdbgui](#) [Download](#) [Docs](#) [Examples](#) [Screenshots](#) [Videos](#) [About](#) [Chat](#) [GitHub](#)

<https://gdbgui.com/>

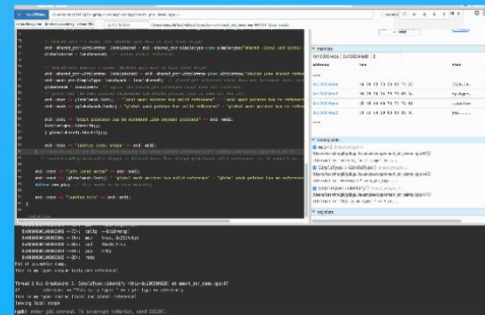
gdbgui

Browser-based debugger for C, C++, go, rust, and more

gdbgui turns this

```
>>> gdb
GNU gdb (Ubuntu 7.11.1-0ubuntu1~16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type 'show copying'
and 'show warranty' for details.
This GDB was configured as 'x86_64-linux-gnu'.
Type 'show configuration' for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type 'help'.
Type 'apropos word' to search for commands related to 'word'.
>>>
```

into this



[Download Now »](#)

[Download Now »](#)

that's all, have fun

"So long
and thanks
for all the fish"