

Luca Tornatore, I.N.A.F.

2025 INAF Course on HPC



Outline

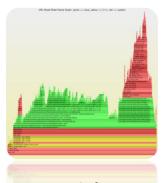




What is perf



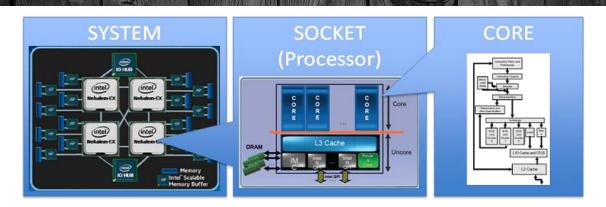
Using perf



Profiling with perf



How well does the CPU work?



The **CPU utilization** number – the figure that the *NIX "top" utility reports – is the **fraction of time slots that** the CPU scheduler in the OS assigned to execution of running programs or the OS itself.

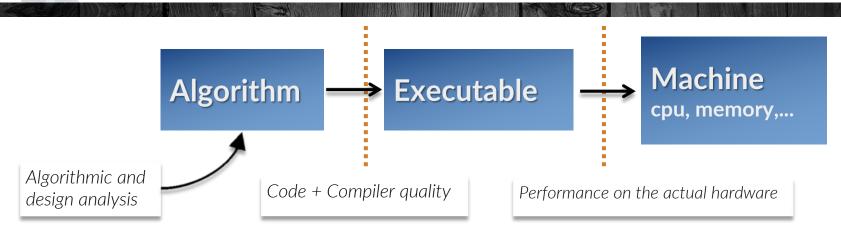
For **compute-bound workloads**, this figure predicted the remaining CPU capacity fairly well for architectures of 80ies and early 90ies.

This metric is now unreliable because of introduction of multi core and multi CPU systems, multi-level caches, non-uniform memory, simultaneous multithreading (SMT), pipelining, out-of-order execution.





Who makes the CPU working well?

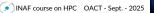


Modern high-end processors provide on-chip hardware that monitors processor's performance and inside events.

That allows to obtain a "precise" and dynamic picture of CPU resources utilization.

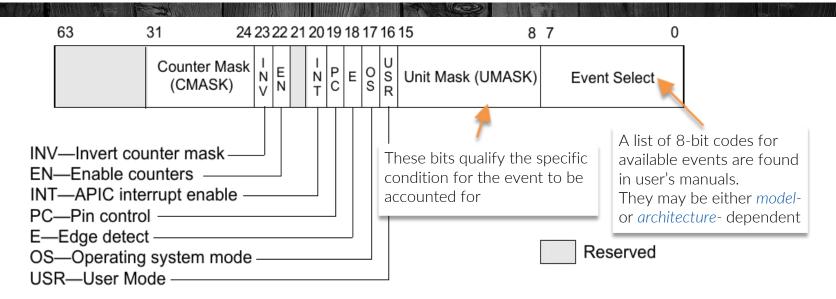
These data can guide performance improvement efforts







Performance Monitoring Units

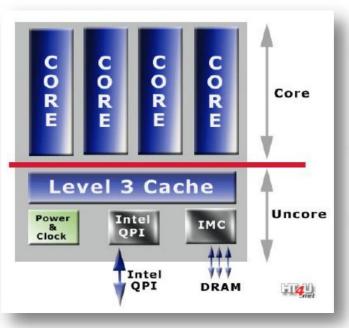


The PMU consists of two types of registers:

- Performance Event Select Register
- Performance Monitor Counter



Type of intecounters



Fixed function counters

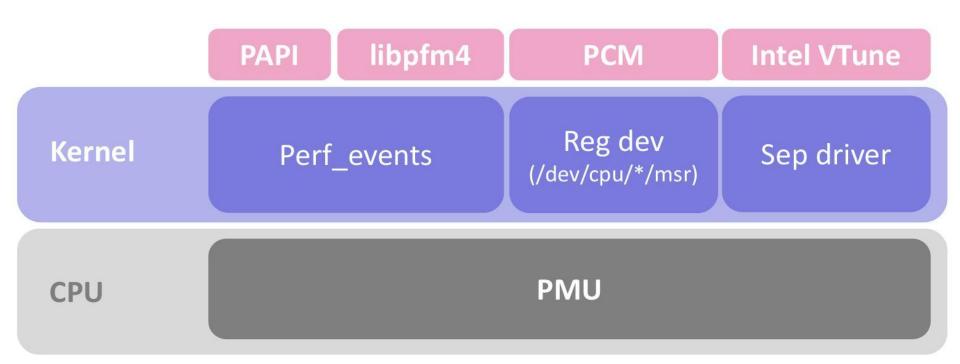
- Predefined events that are commonly used
- TSC, instructions retired, core clock cycles, ...

General purpose performance counters

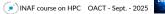
- can be programmed to follow a specific event
- **Precise Event-Based Sampling (PEBS)**
 - Can keep track of architectural state right after instruction causes event
 - Can trigger interrupt (PMI) coupled to counter

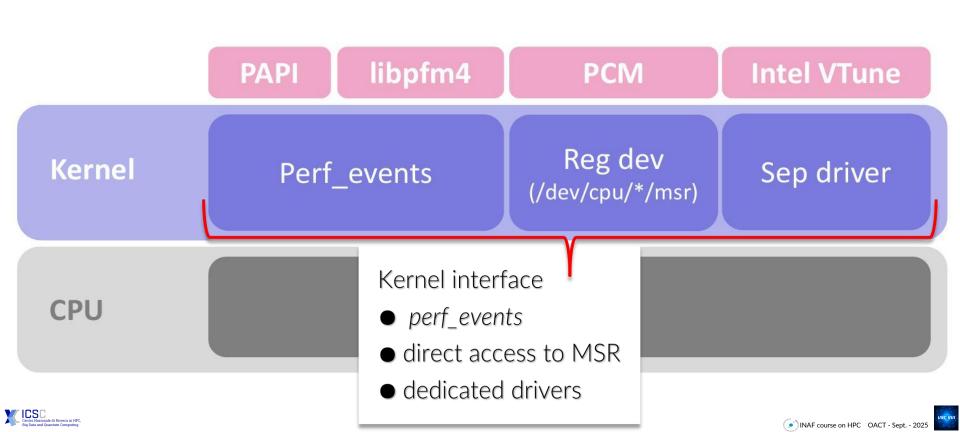


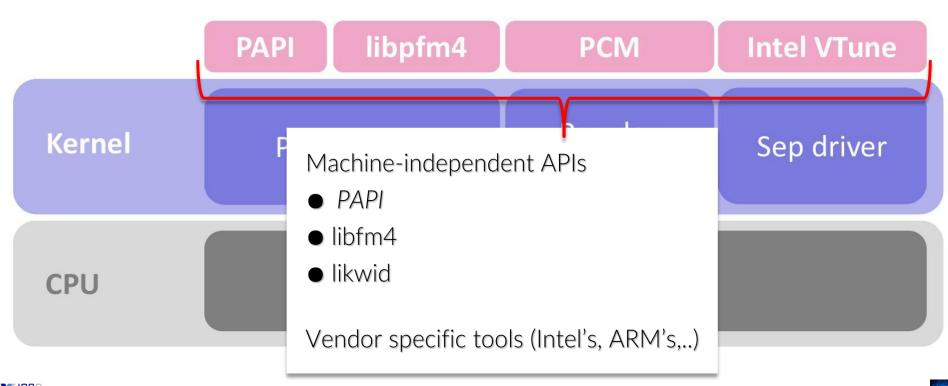
















Direct acces to fixed-function perf. count. through rdpmc

Using PERF, PAPI, PCM, LIKWID, ..., is definitely easier

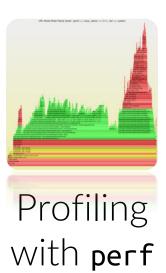
```
01 // rdpmc instructions uses a "fixed-function" performance counter to return the count
           the current core in the low-order 48 bits of an unsigned 64-bit integer.
   unsigned long rdpmc instructions()
04
      unsigned a, d, c;
      c = (1 << 30);
08
       asm volatile("rdpmc" : "=a" (a), "=d" (d) : "c" (c));
10
      return ((unsigned long)a) | (((unsigned long)d) << 32);;</pre>
12
   // rdpmc actual cycles uses a "fixed-function" performance counter to return the coun-
            executed by the current core. Core cycles are not accumulated while the pro-
            which is used when the operating system has no task(s) to run on a processor
   unsigned long rdpmc actual cycles()
18
      unsigned a, d, c;
19
      c = (1 << 30) + 1;
        asm volatile("rdpmc" : "=a" (a), "=d" (d) : "c" (c));
      return ((unsigned long)a) | (((unsigned long)d) << 32);;</pre>
   // rdpmc reference cycles uses a "fixed-function" performance counter to return the co
            CPU core cycles executed by the current core. This counts at the same rate
            when the core is in the "HALT" state. If a timed section of code shows a la
             rdpmc reference cycles, the processor probably spent some time in a HALT sta
   unsigned long rdpmc reference cycles()
      unsigned a, d, c;
      c = (1 << 30) + 2;
        asm volatile("rdpmc" : "=a" (a), "=d" (d) : "c" (c));
36
37
      return ((unsigned long)a) | (((unsigned long)d) << 32);;</pre>
38 }
```

Outline











What is **perf**?

PFRF – Performance Events for Linux – is the standard profiling infrastructure on Linux, built in the kernel since version 2.6.

It strips away the CPUs hardware differences among different systems and presents an abstracted unique command-line interface for performance measurements.

PERF is the standard way to access the hardware performance counters, in both sampling and counting mode.

It consists of:

- a kernel SYSCALL that provide access to both system-software and hardware events:
- a collection of user-space tools to collect, display and analyze performance data.

References:

- https://perf.wiki.kernel.org/index.php/Tutorial
- http://brendangregg.com/perf.html





must-read advice

Using perf requires either to have root privileges or to lower the access protection at kernel level.

That is because perf allows you to spy intimate details of what happens on your machine and as suche may be used for malicious purposes.

The best solution is to enable the users to access the perf events temporarly only when you need it, so that it is not a permanent setting in your kernel.

```
You achieve this goal by opening a shell and typing

sudo echo 0 > /proc/sys/kernel/kptr_restrict

sudo echo -1 > /proc/sys/kernel/perf_event_paranoid
```

A handy solution is to put the two above lines

echo 0 > /proc/sys/kernel/kptr_restrict

echo -1 > /proc/sys/kernel/perf_event_paranoid

into a text file, for instance named enable_perf in any path path/to/script you prefer, making it executable with mod +x enable perf and then execute it from a shell as sudo

sudo path/to/script/enable_perf



```
What is perf?
```

Hello perf

```
luca@GGG:~% perf
usage: perf [--version] [--help] [OPTIONS] COMMAND [ARGS]
The most commonly used perf commands are:
                   Read perf.data (created by perf record) and display annotated code
   annotate
   archive
                   Create archive with object files with build-ids found in perf.data file
                   General framework for benchmark suites
   bench
   buildid-cache
                   Manage build-id cache.
   buildid-list
                   List the buildids in a perf.data file
  c2c
                   Shared Data C2C/HITM Analyzer.
                   Get and set variables in a configuration file.
   config
   data
                   Data file related processing
   diff
                   Read perf.data files and display the differential profile
   evlist
                   List the event names in a perf.data file
  ftrace
                   simple wrapper for kernel's ftrace functionality
  inject
                   Filter to augment the events stream with additional information
   kallsyms
                   Searches running kernel for symbols
                   Tool to trace/measure kernel memory properties
   kmem
                   Tool to trace/measure kvm quest os
   kvm
   list
                   List all symbolic event types
   lock
                   Analyze lock events
                   Profile memory accesses
   mem
                   Run a command and record its profile into perf.data
   record
                   Read perf.data (created by perf record) and display the profile
   report
   sched
                   Tool to trace/measure scheduler properties (latencies)
   script
                   Read perf.data (created by perf record) and display trace output
   stat
                   Run a command and gather performance counter statistics
   test
                   Runs sanity tests.
   timechart
                   Tool to visualize total system behavior during a workload
                   System profiling tool.
   top
   probe
                   Define new dynamic tracepoints
  trace
                   strace inspired tool
```



What is **perf**?

Hello perf

luca@GGG:~% perf

usage: perf [--version] [--help] [OPTIONS] COMMAND [ARGS]

The most commonly used perf commands are:

annotate archive bench buildid-cache buildid-list c2c config data diff evlist ftrace inject kallsyms kmem kvm list lock mem record

record report sched script stat

test timechart top probe

trace

Read perf.data (created by perf record) and display annotated code Create archive with object files with build-ids found in perf.data file General framework for benchmark suites Manage build-id cache.

List the buildids in a perf.data file Shared Data C2C/HITM Analyzer.

Get and set variables in a configuration file.

Data file related processing

Read perf.data files and display the differential profile List the event names in a perf.data file

simple wrapper for kernel's ftrace functionality

Filter to augment the events stream with additional information

Searches running kernel for symbols

Tool to trace/measure kernel memory properties

Tool to trace/measure kvm guest os

List all symbolic event types Analyze lock events

Profile memory accesses

Run a command and record its profile into perf.data Read perf.data (created by perf record) and display the profile

Tool to trace/measure scheduler properties (latencies)

Read perf.data (created by perf record) and display trace output

Run a command and gather performance counter statistics Runs sanity tests.

Tool to visualize total system behavior during a workload System profiling tool.

Define new dynamic tracepoints

Define new dynamic tracepoints strace inspired tool

list of options for each command use perf command -h

Note: to obtain the



TYPE OF EVENTS

The kernel interface can measure and report events from several different sources:

- software events: pure software events (ex: context switches)
- PMU hardware events: Performance Monitoring Unit (PMU) microarchitectural events (ex. Nr. of instructions retired, branch-misses, ...)
- hardware events: common hardware events with mnemonic names. mapped on actual CPU events
- tracepoint events: implemented through the kernel ftrace infrastracture.





What is

perf?

HOW MANY EVENTS IN YOUR CPU?

Let's have a look:

- perf list
- showevtinfo [pmu-tools]
- papi_avail [PAPI]
- ▶ likwid-perfctr -e
- https://download.01.org/perfmon/

...hundreds of different events



Names for **perf** Events

Some commonly used events have mnemonic names:

List of pre-defined events (to be used in -e): branch-instructions OR branches [Hardware event] branch-misses [Hardware event] bus-cycles [Hardware event] cache-misses [Hardware event] cache-references [Hardware event] cpu-cycles OR cycles [Hardware event] instructions [Hardware event] ref-cycles [Hardware event] alignment-faults [Software event] bpf-output [Software event] context-switches OR cs [Software event] cpu-clock [Software event] cpu-migrations OR migrations [Software event] dummy [Software event] emulation-faults [Software event] major-faults [Software event] minor-faults [Software event] page-faults OR faults [Software event] task-clock [Software event] L1-dcache-load-misses [Hardware cache event] L1-dcache-loads [Hardware cache event] L1-dcache-stores [Hardware cache event] L1-icache-load-misses [Hardware cache event] LLC-load-misses [Hardware cache event] LLC-loads [Hardware cache event] LLC-store-misses [Hardware cache event] LLC-stores [Hardware cache event] branch-load-misses Hardware cache event branch-loads [Hardware cache event] dTLB-load-misses [Hardware cache event dTLB-loads [Hardware cache event]



perf Events

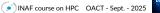
An event can also be specified by its processor-specific identifier, that usually is a hex value

Example:

```
"EventCode": "0x00",
    "UMask": "0x01",
    "EventName": "INST_RETIRED.ANY",
    "BriefDescription": "Instructions retired from execution."

"EventCode": "0x00",
    "UMask": "0x02",
    "EventName": "CPU_CLK_UNHALTED.THREAD",
    "BriefDescription": "Core cycles when the thread is not in halt state"
```





You can specify events by their raw code if not present in **perf list**perf stat -e r5100c0 ...

You may obtain the hex code in several way

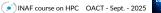
1. Look in the list provided by Intel, combining "Umask" and "Event Code" fields

```
perf stat -e r<umask><eventselector> ...
perf stat -e cpu/event=0xcode,umask=0xcode/u ...
```

1. Use libpfm4, using showevtinfo and then check_events:

```
check events <event name>:<umask>[(:modifers)*]
```



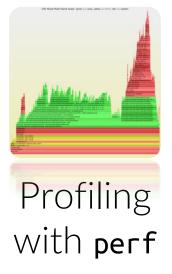


Outline











Right answers to right questions

- Which code segments take most of the execution time? What are the causes? (cache misses, branch misses, non-optimal pipelines..)
- Do the performance counters suggest something about performance issues?
- Do the performance counters offer some hints on how to fix the issues?

Workflow:

- Run the program and collect as many data as possible
- Analyse the collected data and profile the program





Right answers to right questions





EVENT-based profiling

can be accessed from inside the code collecting precise values of events counters in a given time range

SAMPLE-based profiling

periodic sampling of program/system status and of event counters values









EVENT-based profiling

can be accessed from inside the code collecting precise values of events counters in a given time range





Counting events with **Perf**

For all supported events, perf stat can count during a process execution:

```
luca@GGG:~% perf stat -B dd if=/dev/zero of=/dev/null count=1000000
1000000+0 records in
1000000+0 records out
512000000 bytes (512 MB, 488 MiB) copied, 1.0796 s, 474 MB/s
 Performance counter stats for 'dd if=/dev/zero of=/dev/null count=1000000':
                                                        0.997 CPUs utilized
       1078.764601
                        task-clock (msec)
                        context-switches
                                                        0.008 K/sec
                        cpu-migrations
                                                        0.001 K/sec
                74
                        page-faults
                                                        0.069 K/sec
     2,152,108,887
                        cycles
                                                        1.995 GHz
     1,986,849,204
                        instructions
                                                        0.92 insn per cycle
       387,532,518
                        branches
                                                     359.237 M/sec
         8,022,754
                        branch-misses
                                                        2.07% of all branches
       1.081948899 seconds time elapsed
```







Counting events with **Perf**

You can specify the events you want to profile by their mnemonic names:

```
perf stat -e cycles:.,instructions:.,branch-misses:.,cache-misses:.
```

If you specify more events than available counters, the kernel uses time multiplexing to sample all the events; then perf scales the count based on running time vs the total amount of time that a given event has been active

```
final count = raw count * time enabled / time runnig
```







Counting events with **Perf**

Processor-wide mode:

```
perf stat -e cycles:u -a ./executable arguments
```

Sampling a subset of cores:

```
perf stat -e cycles:u -a -C 1,4-5 ./executable arguments
```

Profile a given process:

```
perf stat -e cycles:u -p PID sleep 2
```

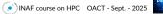
Pretty printing of large numbers

```
perf stat -B -e cycles:u -e instructions:u -e branch-misses:u ./executable arguments
```

Repeat and average measures

```
perf stat -r 10 -B -e cycles:u -e instructions:u ./executable arguments
```





Right answers to right questions



SAMPLE-based data collection

periodic sampling of program/system status and of event counters values







Collecting information with **Perf**

You can collect much more information, with line-level detail:

```
perf record -e cycles:.,instructions:.,branch-misses:.,cache-misses:.
```

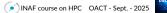
A file named perf.data will be produced, with all the data collected during the perf activity.

At odds with **perf stat**, this mode profiles CPU usage based on sampling the instruction pointer or stack tree at a given rate:

```
perf record -F fff -a -k CLOCK MONOTONIC -- sleep 10
```

Sample the whole systems (-a, all CPUs) at fff Hertz (-F fff) without accounting for the perf activity (--) for 10 secs.









Sampling the call stack

You can collect dynamic information about the call-stack (so that to obtain the call graph)

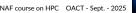
```
perf record --call-graph <record mode [,record size ]>
```

Where record_mode can be <fp|dwarf|lbr> and record_size is the max size of stack recording (in bytes) for dwarf mode (def: 8192)

```
You can specify -a (all CPUs); -C, --cpu <cpu>; -s, --per-thread;
```

Now, let's go back to our friend lie.c to check whether perf is mislead as much as gprof was...









Sampling the call stack

perf report --call-graph graph, 0.6, 100, caller --stdio --no-children

```
Total Lost Samples: 0
Samples: 740 of event (cycles:ppp)
Event count (approx.): 556806058
Overhead Command Shared Object
                                    Symbol
  99.61% lie1.q lie1.q
                              [.] loop
          --- start
              libc start main
            main
             --98.73%--heavy
                       loop
             --0.88%--light
                       loop
```

```
# Total Lost Samples: 0
 Samples: 740 of event cycles:ppp
 Event count (approx.): 556806058
# Overhead Command Shared Object
                                     Symbol
   99.61% lie1.g lie1.g
                                     [.] loop
           ---loop
               --98.73% -- heavy
                        main
                          libc start main
                        start
               --0.88%--light
                        main
                          libc start main
```

caller order

callee order







Collecting information with perf

As well, you can record performance events along with the call stack

```
perf record --call-graph ... -e event_1:.,event_2:.,event_3:. ...
```

With all the same options valid before.

Check perf record -h for more details.

Note:

remember to compile (with gcc) with
-fno-omit-frame-pointer
in order to reconstruct the call stack





Outline









Recovering the call stack

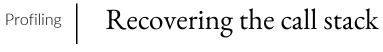
Once you have perf.data file, you can easily get the call graph

perf report -call-graph --stdio

```
# Samples: 52K of event 'cycles:.pppu'
# Event count (approx.): 32241477547
# Children
               Self Command
                                      Shared Object
                                                          Symbol
  100.00%
              0.00% prefetching on. prefetching on.03 [.] start
            --- start
                libc start main
              main
               --0.68% -- rand
                           random
              0.00% prefetching on. libc-2.27.so
  100.00%
                                                          [.] libc start main
            --- libc start main
              main
               --0.68% -- rand
                           random
             98.82% prefetching on. prefetching on.03 [.] main
  100.00%
             --98.82%-- start
                        libc start main
                      main
            --1.18%--main
                       --0.68%--rand
                                 __random
              0.08% prefetching on. libc-2.27.so
    0.72%
                                                          [.] rand
            --0.64%--rand
                      __random
    0.68%
              0.31% prefetching on. libc-2.27.so
                                                          [.] random
    0.43%
              0.43% prefetching on. [unknown]
                                                          [k] 0xfffffffff8e0009e7
     0.33%
              0.33% prefetching on. libc-2.27.so
                                                          [.] random r
```



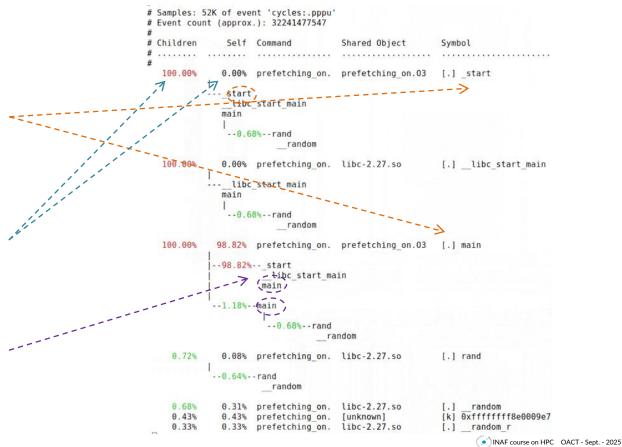




The symbol this call refers to

Children time and self time are reported separately

Both caller and callee informations







Recovering the call stack

```
perf report -call-graph -stdio \
--no-children
```

```
# Samples: 52K of event 'cycles:.pppu'
# Event count (approx.): 32241477547
# Overhead Command
                            Shared Object
                                               Symbol
   98.82% prefetching on. prefetching on.03 [.] main
            --98.81%--main
                        libc start main
                       start
    0.43%
           prefetching on.
                           [unknown]
                                               [k] 0xffffffff8e0009e7
    0.33%
           prefetching on.
                            libc-2.27.so
                                                   random r
    0.31%
           prefetching on.
                           libc-2.27.so
                                                    random
    0.08%
           prefetching on.
                            libc-2.27.so
                                                   rand
    0.03%
           prefetching on.
                            prefetching on.03
                                                   rand@plt
    0.00%
           prefetching on.
                            libc-2.27.so
                                                   dl addr
    0.00%
           prefetching on.
                            ld-2.27.so
                                               [.] dl relocate object
    0.00%
           prefetching on.
                            ld-2.27.so
                                               [.] do lookup x
    0.00%
           prefetching on.
                            ld-2.27.so
                                               [.] strcmp
    0.00%
           prefetching on.
                            ld-2.27.so
                                               [.] GI tunables init
```





perf report -call-graph -stdio \ --no-children

callee format as default

```
# Samples: 52K of event 'cycles:.pppu'
# Event count (approx.): 32241477547
# Overhead Command
                            Shared Object
                                               Symbol
   98.82% prefetching on. prefetching on.03 [.] main
             --98.81%--main
                        libc start main
                       start
    0.43%
           prefetching on.
                            [unknown]
                                                [k] 0xfffffffff8e0009e7
                            libc-2.27.so
           prefetching on.
                                                    random r
    0.31%
           prefetching on.
                            libc-2.27.so
                                                    random
    0.08%
           prefetching on.
                            libc-2.27.so
                                                   rand
    0.03%
           prefetching on.
                            prefetching on.03
                                                   rand@plt
    0.00%
           prefetching on.
                            libc-2.27.so
                                                   dl addr
    0.00%
           prefetching on.
                            ld-2.27.so
                                                   dl relocate object
    0.00%
           prefetching on.
                            ld-2.27.so
                                                [.] do lookup x
    0.00%
           prefetching on.
                            ld-2.27.so
                                                [.] strcmp
    0.00%
           prefetching on.
                            ld-2.27.so
                                                [.] GI tunables init
```





```
perf report --call-graph --stdio \
--no-children --sort=dso
```

Different ordering are possible: comm, dso, symbol,...

```
# Samples: 52K of event 'cycles:.pppu'
# Event count (approx.): 32241477547
# Overhead Shared Object
   98.84% prefetching on.03
            --- start
                 libc start main
               main
    0.72% libc-2.27.so
            --- start
                 libc start main
               main
                --0.68%--rand
                            random
    0.43% [unknown]
```



```
perf report --call-graph --stdio \
--no-children --sort=comm
```

Different ordering are possible: comm, dso, symbol,...

```
# Samples: 52K of event 'cycles:.pppu'
# Event count (approx.): 32241477547
 Overhead Command
   100.00% prefetching on.
             --100.00%-- start
                         libc start main
                       main
                        --0.68%--rand
                                    random
```





```
perf report --call-graph --stdio \
--no-children
```

The graph for all the events recorderd are shown

```
# Samples: 55K of event (instructions:.u')
# Event count (approx.): 9679409860
# Overhead Command
                            Shared Object
                                               Symbol
   97.01% prefetching on. prefetching on.03 [.] main
           --- start
               libc start main
              main
    2.05% prefetching on. libc-2.27.so
                                               [.] random
            --- start
               libc start main
              main
              rand
                random
```





Detailed browsing of the events

You can browse your code at line-level detail, with reported metrics for all the events that you requested

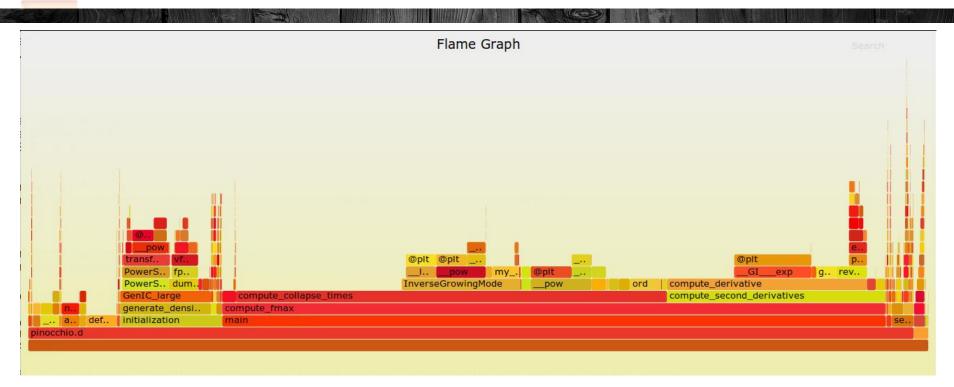
 \rightarrow see live examples

```
uca/code/HPC LECTURES/prefetching/prefetching on.g03
                     $0x1f,%eax
 0.09
              add
                     %edx, %eax
 0.06
                     %eax
                  mid = (low + high) / 2;
 8.90
                     %eax, %edx
               while(low <= high) {
 0.01
                     %esi,%ecx
            1 jq
                   builtin prefetch (&data[(low + mid - 1)/2], 0, 1);
15.96
        1f7:
                     -0x1(%rcx,%rdx,1),%r10d
 0.13
                     %r10d, %eax
 1.06
                     $0x1f,%eax
              shr
 0.50
                     %r10d,%eax
                   builtin prefetch (&data[(mid + 1 + high)/2], 0, 1);
 3.68
                    0x1(%rdx),%r10d
                   builtin prefetch (&data[(low + mid - 1)/2], 0, 1);
 0.10
                   builtin prefetch (&data[(mid + 1 + high)/2], 0, 1);
 8.42
                     (%r10,%rsi,1),%r11d
                   builtin prefetch (&data[(low + mid - 1)/2], 0, 1);
 0.12
              clta
 0.64
              prefetcht0 (%r14,%rax,4)
                   builtin prefetch (&data[(mid + 1 + high)/2], 0, 1);
23.31
                     %rlld,%eax
 0.08
              shr
                     $0x1f,%eax
 0.11
                     %rlld,%eax
 0.11
              sar
                     %eax
 1.64
              movslq %eax, %r11
 0.11
              prefetcht0 (%r14,%r11,4)
                 if(data[mid] < Key)
 3.92
              movslq %edx,%r11
 0.03
              cmp
                     (%r14,%r11,4),%edi
 0.05
            † jle
                     le0
                   low = mid + 1:
 3.61
                     %r10d.%ecx
                  mid = (low + high) / 2;
 0.08
                     %eax, %edx
               while(low <= high) {
 0.08
                     %esi,%ecx
 0.01
            † jle
                     1f7
 1.98
       23b:
              add
                     $0x4,%r15
            main():
              for (i = 0; i < Nsearch; i++)
 1.75
                     %r12,%r15
            † ine
                     1d0
                  found++;
Press 'h' for help on key bindings
```



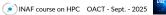
Profiling

Flame graphs

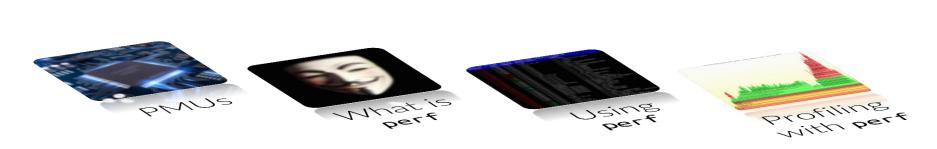


Have a look at Brend Gregg's blog: http://www.brendangregg.com/FlameGraphs/cpuflamegraphs.html





Outline



Addendum I

Elementary intro to PAPI







(very) Basic usage of PAPI

INITIAL IZATION

- define events you want to profile
- reserve room for counter values
- initialize library
- program each event separately

```
#define PCHECK(e) \
 if (e!=PAPI OK)
    {printf("Problem in papi call, line %d\n", LINE ); return 1;}
#define NEVENTS 3
int main(int argc, char **argv)
 int events[NEVENTS] =
     PAPI TOT CYC,/* total cycles */
     PAPI_L1_DCM, /* stalls on L1 cache miss */
     PAPI_L2_DCM, /* stalls on L2 cache miss */
  long long
                        values[NEVENTS];
                        retval;
 retval = PAPI library init(PAPI VER CURRENT);
 if (retval != PAPI VER CURRENT)
   printf("wrong PAPI initialization: %d instead of %d\n", retval, PAPI VER CURRENT);
    for ( int i = 0; i < NEVENTS; i++ )</pre>
        retval = PAPI query event(events[i]);
       PCHECK(retval);
```



(very) Basic usage of PAPI

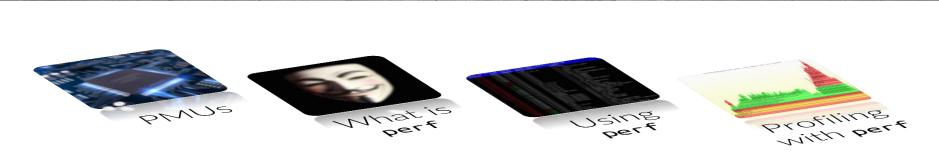
```
retval = PAPI start counters(events, NEVENTS); PCHECK(retval);
/* run the experiment */
for (i=0; i<NRUNS; i++)</pre>
    for (j=0; j<size; j++) array[j] = 2.3*array[j]+1.2;
retval = PAPI stop counters(values, NEVENTS); PCHECK(retval);
printf("size: %d cycles: %lld cycles_per_loc: %9.5f L1miss: %lld <...>
       size.
       values[0], values[0]/(1.*NRUNS*size),
       values[1], (double)values[1]/(NRUNS*size),
       values[2], (double)values[2]/(NRUNS*size));
```

Start, stop and access





Outline



Addendum II

Using gperftools



Use gperftools to obtain a basic profile

OPTION I

Include gperftools/profiler.h:

#include <gperftools/profiler.h>

- Encompass the code segments to be profiled within calls:

```
ProfilerStart ( "name_of_profile_file" );
...
ProfilerStop ( );
```

- Then variable CPUPROFILEFREQUENCY=x modifies the sampling frequency (x is in Hz)

OPTION II

- link exec. against libprofiler.so

cc source.c -o exec -lprofiler

- or pre-load the profiler libraryLD_PRELOAD=/path/to/libprofiler.so ./exec
- set variables CPUPROFILE to start the profiling

LD_PRELOAD=/path/to/libprofiler.so \
CPUPROFILE=./exec.prof CPUFREQUENCY=1000 \
./exec parameters





Use gperftools to obtain a basic profile

READ THE GPERFTOOLS DATA

Use pprof, interactive mode:

pprof ./exec ./exec.prof

... let's have a live demonstration about:

top

- list
- disasm
- weblist
- gv / web

You can produce output on stdio:

```
pprof --text ./exec ./exec.prof
pprof --text --lines ./exec ./exec.prof
pprof --text --functions ./exec ./exec.prof
```

```
or in other image format pprof <--gv|--web|--dot|--ps|--svg|--gif>
```

or output to be pipelined elsewhere pprof <--raw|--collapsed|--callgrind>

or the whole call stack in detail pprof --text --stack

